

---

# EDDI: Efficient Dynamic Discovery of High-Value Information with Partial VAE

---

Chao Ma<sup>1</sup> Sebastian Tschitschek<sup>2</sup> Konstantina Palla<sup>2</sup> José Miguel Hernández-Lobato<sup>1 2</sup>  
Sebastian Nowozin<sup>3</sup> Cheng Zhang<sup>2</sup>

## Abstract

Many real-life decision making situations allow further relevant information to be acquired at a specific cost, for example, in assessing the health status of a patient we may decide to take additional measurements such as diagnostic tests or imaging scans before making a final assessment. Acquiring more relevant information enables better decision making, but may be costly. How can we trade off the desire to make good decisions by acquiring further information with the cost of performing that acquisition? To this end, we propose a principled framework, named *EDDI* (Efficient Dynamic Discovery of high-value Information), based on the theory of Bayesian experimental design. In *EDDI*, we propose a novel *partial variational autoencoder* (Partial VAE) to predict missing data entries problematically given any subset of the observed ones, and combine it with an acquisition function that maximizes expected information gain on a set of target variables. We show cost reduction at the same decision quality and improved decision quality at the same cost in multiple machine learning benchmarks and two real-world health-care applications.

## 1 Introduction

Imagine a person walking into a hospital with a broken arm. The first question from health-care personnel would likely be “How did you break your arm?” instead of “Do you have a cold?”, because the answer reveals relevant information

for this patient’s treatment. Human experts dynamically acquire information based on the current understanding of the situation. Automating this human expertise of asking relevant questions is difficult. In other applications such as online questionnaires for example, most existing online questionnaire systems either present exhaustive questions (Lewenberg et al., 2017; Shim et al., 2018) or use extremely time-consuming human labeling work to manually build a decision tree to reduce the number of questions (Zakim et al., 2008). This wastes the valuable time of experts or users (patients). An automated solution for personalized dynamic acquisition of information has great potential to save much of this time in many real-life applications.

What are the technical challenges to building an intelligent information acquisition system? *Missing data is a key issue*: taking the questionnaire scenario as an example, at any point in time we only observe a small subset of answers yet have to reason about possible answers for the remaining questions. We thus need an accurate probabilistic model that can perform inference given a variable subset of observed answers. *Another key problem is deciding what to ask next*: this requires assessing the value of each possible question or measurement, the exact computation of which is intractable. However, compared to current active learning methods we select individual features, not instances; therefore, existing methods are not applicable. In addition, these traditional methods are often not scalable to the large volume of data available in many practical cases (Settles, 2012; Lewenberg et al., 2017).

We propose the *EDDI* (Efficient Dynamic Discovery of high-value Information) framework as a scalable information acquisition system for any given task. We assume that information acquisition is always associated with some cost. Given a task, such as estimating the customers’ experience or assessing population health status, we dynamically decide which piece of information to acquire next. The framework is very general, and the information can be presented in any form such as answers to questions, or results of lab tests. Our contributions are:

---

<sup>1</sup>Department of Engineering, University of Cambridge, Cambridge, UK <sup>2</sup>Microsoft Research, Cambridge, UK <sup>3</sup>Google AI, Berlin, Germany (contributed while being with Microsoft Research). Correspondence to: Cheng Zhang <Cheng.Zhang@microsoft.com>.

- We propose a novel efficient information acquisition framework, EDDI (Section 3). To enable EDDI, we contribute technically:

1. *A new partial amortized inference method for generative modeling under partially observed data (Section 3.2).* We extend the variational autoencoder (VAE) (Kingma & Welling, 2014; Rezende et al., 2014), to account for partial observations. The resulting method, which we call the Partial VAE, is inspired by the set formulation of the data (Qi et al., 2017; Zaheer et al., 2017). The Partial VAE, as a probabilistic framework in the presence of missing data, is highly scalable, and serves as the base for the EDDI framework. Note that Partial VAE itself is widely applicable and can be used on its own as a non-linear probabilistic framework for missing-data imputation.
2. *An information theoretic acquisition function with a novel efficient approximation, yielding a novel variable-wise active learning method (Section 3.3).*

Based on the partial VAE, we actively select the unobserved variable which contributes most to the task, such as costumer surveys and health assessments, evaluated using the mutual information. This acquisition function does not have an analytical solution, and we derive a novel efficient approximation.

- We demonstrate the performance of EDDI in various settings, and apply it in real-life health-care scenarios (Section 4).
1. We first show the superior performance of the Partial VAE framework on an image inpainting task (Section 4.1).
  2. We then use 6 different datasets from the Machine Learning repository of University of Irvine (UCI) (Dheeru & Karra Taniskidou, 2017) to demonstrate the behavior of EDDI, comparing with multiple baseline methods (Section 4.2).
  3. Finally, we evaluate EDDI on two real-life health-care applications: risk assessment in intensive care (Section 4.3) and public health assessment using a national survey (Section 4.4), where traditional methods without amortized inference do not scale. EDDI shows clear improvements in both applications.

## 2 Related Work

EDDI requires a method that handles partially observed data to enable dynamic variable wise active learning. We thus review related methods for handling partial observation and performing active learning.

### 2.1 Partial Observation

Missing data entries are common in many real-life applications, which has created a long history of research on the topic of dealing with missing data (Rubin, 1976; Dempster et al., 1977). We describe existing methods below with the focus of probabilistic methods:

**Traditional methods without amortization.** Prediction based methods have shown advantages for missing value imputation (Scheffer, 2002). Efficient matrix factorization based methods have been recently applied (Keshavan et al., 2010; Jain et al., 2010; Salakhutdinov & Mnih, 2008), where the observations are assumed to be able to decompose as the multiplication of low dimensional matrices. In particular, many probabilistic frameworks with various distribution assumptions (Salakhutdinov & Mnih, 2008; Blei et al., 2003) have been used for missing value imputation (Yu et al., 2016; Hamesse et al., 2018) and also recommender systems where unlabeled items are predicted (Stern et al., 2009; Wang & Blei, 2011; Gopalan et al., 2014).

The probabilistic matrix factorization method has been used in the active variable selection framework called the dimensionality reduction active learning model (DRAL), (Lewenberg et al., 2017). These traditional methods suffer from limited model capacity since they are typically linear. Additionally, they do not scale to large volumes of data and thus are usually not applicable in real-world applications. For example, Lewenberg et al. (2017) tested the performance of their method with a single user due to the heavy computational cost of traditional inference methods for probabilistic matrix factorization.

**Utilizing Amortized Inference.** Amortized inference (Kingma & Welling, 2014; Rezende et al., 2014; Zhang et al., 2017) has significantly improved the scalability of deep generative latent variable models. In the case of partially observed data, amortized inference is particularly of interest due to the speed requirement in many real-life applications. Wu et al. (2018) use amortized inference during training, where the training dataset is assumed to be fully observed. During test time, the traditional non-scalable inference is used to infer missing data entries from the partially observed dataset using the pre-trained model. This method is restrictive since it is not scalable in the test time and the fully observed training set assumption does not hold for many applications.

Nazabal et al. (2018) use zero imputation (ZI) for amortized inference for both training and test sets with missing data entries. ZI is a generic and straightforward method that first fills the missing data with zeros, and then feeds the imputed data as input for the inference network. The drawback of ZI is that it introduces bias when the data are not missing completely at random which leads to a poorly

fit model. We also observe artifacts when using it for the image inpainting task. Independent of our work, [Garnelo et al. \(2018\)](#) explore interpreting variational autoencoder (amortized inference) as stochastic processes, which also handles partial observation per se.

## 2.2 Active Learning

**Traditional Active Learning.** Active learning, also referred to as experimental design, aims to obtain optimal performance with fewer selected data (or experiments) ([Lindley, 1956](#); [MacKay, 1992](#); [Settles, 2012](#)). Traditional active learning aims to select the *next data point* to label. Many information theoretical approaches have shown promising results in various settings with different acquisition functions ([MacKay, 1992](#); [McCallumzy & Nigamy, 1998](#); [Houlsby et al., 2011](#)). These methods commonly assume that the data are fully observed, and the acquisition decision is instance wise. Little work has dealt with missing values within instances. [Zheng & Padmanabhan \(2002\)](#) deal with missing data values by imputing with traditional non-probabilistic methods ([Little & Rubin, 1987](#)) first. It is still an instance-wise active learning framework.

Different from traditional active learning, our proposed framework performs *variable-wise active learning for each instance*. In this setting, information theoretical acquisition functions need a new design as well as non-trivial approximations. The most closely related work is the aforementioned DRAL ([Lewenberg et al., 2017](#)), which deals with variable-wise active learning for each instance.

**Active Feature Acquisition (AFA).** Active sequential feature selection is of great need, especially in cost-sensitive applications. Thus, many methods have also been applied and resulted in the class of methodologies called Active Feature Acquisition (AFA) ([Melville et al., 2004](#); [Saar-Tsechansky et al., 2009](#); [Thahir et al., 2012](#); [Huang et al., 2018](#)). For instance, [Melville et al. \(2004\)](#); [Saar-Tsechansky et al. \(2009\)](#) have designed objectives to select any feature from any instance to minimize the cost to achieve high accuracy. The proposed framework is very general. However, the problem setting of AFA methods is different from our active variable selection problem. AFA aims to select training set optimally that would result in the best classifier (model), while assume that the test data are fully observed. On the contrary, our framework aims to identify and acquire high value information sequentially for each test instance.

## 3 Method

In this section, we first formalize the active variable selection problem. Then, we present the Partial VAE to model and perform inference on partial observations. Finally, we complete the EDDI framework by presenting our new acquisition function and estimation method.

### 3.1 Problem formulation

In this work, we focus on the following active variable selection problem. Let  $\mathbf{x} = [x_1, \dots, x_{|I|}]$  be a set of random variables with probability density  $p(\mathbf{x})$ . Furthermore, let a subset of the variables  $\mathbf{x}_O$ ,  $O \subset I$ , be observed while the variables  $\mathbf{x}_U$ ,  $U = I \setminus O$ , are unobserved. Assume that we can query the value of variables  $x_i$  for  $i \in U$ . The goal of active variable selection is to query a sequence of variables in  $U$  in order to predict a quantity of interest  $f(\mathbf{x})$ , as accurately as possible while simultaneously performing as few queries as possible, where  $f(\cdot)$  can be any (random) function. This problem, in the simplified myopic setting, can be formalized as that of proposing the next variable  $x_{i^*}$  to be queried by maximizing a reward function  $R$  at each step:

$$i^* = \arg \max_{i \in U} R(i | \mathbf{x}_O), \quad (1)$$

where  $R(i | \mathbf{x}_O)$  quantifies the merit of our prediction of  $f(\cdot)$  given  $\mathbf{x}_O$  and  $x_i$ . Furthermore, the reward can quantify other properties important to the problem, e.g. the cost of acquiring  $x_i$ .

### 3.2 Partial Amortization of Inference Queries

We first introduce how to establish a generative probabilistic model of random variables  $\mathbf{x}$ , that is capable of handling unobserved (missing) variables  $\mathbf{x}_U$  with variable size. Our approach to this, named the Partial VAE, is based on the variational autoencoder (VAE), which enables amortized inference to scale to large volumes of data.

**VAE and amortized inference.** A VAE defines a generative model in which the data  $\mathbf{x}$  is generated from latent variables  $\mathbf{z}$ ,  $p(\mathbf{x}, \mathbf{z}; \theta) = \prod_i p_\theta(\mathbf{x}_i | \mathbf{z}) p(\mathbf{z})$ . The data generation,  $p_\theta(\mathbf{x} | \mathbf{z})$ , is realized by a deep neural network. To approximate the posterior of the latent variable  $p_\theta(\mathbf{z} | \mathbf{x})$ , VAEs use *amortized* variational inference. Specifically, it uses an encoder, which is another neural network with the data  $\mathbf{x}$  as input to produce a variational approximation of the posterior  $q(\mathbf{z} | \mathbf{x}; \phi)$ . As traditional variational inference, VAE is trained by maximizing an evidence lower bound (ELBO), which is equivalent to minimizing the KL divergence between  $q(\mathbf{z} | \mathbf{x}; \phi)$  and  $p_\theta(\mathbf{z} | \mathbf{x})$ .

VAEs are not directly applicable when data points have arbitrary subset of data entries missing. Consider the situation that the variables are divided into *observed* variables  $\mathbf{x}_O$  and *unobserved* variables  $\mathbf{x}_U$ . In this setting, we would like to efficiently and accurately infer  $p(\mathbf{z} | \mathbf{x}_O)$  and  $p(\mathbf{x}_U | \mathbf{x}_O)$ . One main challenge is that there are many possible partitions  $\{U, O\}$ , where the size of observed variables might vary. Therefore, classic approaches to training a VAE with the variational bound and amortized inference networks are not applicable. We propose to extend amortized inference to handle partial observations.

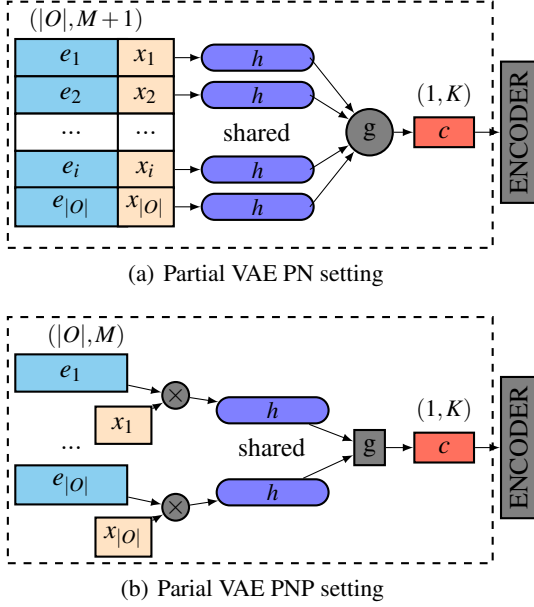


Figure 1: Illustration of Partial VAE encoder architecture.

**Partial VAE.** In a VAE,  $p(\mathbf{x}|\mathbf{z})$  is factorized, i.e.

$$p(\mathbf{x}|\mathbf{z}) = \prod_i p_i(\mathbf{x}_i|\mathbf{z}). \quad (2)$$

This implies that given  $\mathbf{z}$ , the observed variables  $\mathbf{x}_O$  are conditionally independent of  $\mathbf{x}_U$ . Therefore,

$$p(\mathbf{x}_U|\mathbf{x}_O, \mathbf{z}) = p(\mathbf{x}_U|\mathbf{z}), \quad (3)$$

and inferences about  $\mathbf{x}_U$  can be reduced to inference about  $\mathbf{z}$ . Hence, the key object of interest in this setting is  $p(\mathbf{z}|\mathbf{x}_O)$ , i.e., the posterior over the latent variables  $\mathbf{z}$  given the observed variables  $\mathbf{x}_O$ . Once we obtain  $\mathbf{z}$ , computing  $\mathbf{x}_U$  is straightforward. To approximate  $p(\mathbf{z}|\mathbf{x}_O)$ , we introduce a variational inference network  $q(\mathbf{z}|\mathbf{x}_O)$  and define a partial variational lower bound,

$$\begin{aligned} \log p(\mathbf{x}_O) &\geq \log p(\mathbf{x}_O) - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}_O) \| p(\mathbf{z}|\mathbf{x}_O)) \quad (4) \\ &= \mathbb{E}_{\mathbf{z} \sim q(\mathbf{z}|\mathbf{x}_O)} [\log p(\mathbf{x}_O|\mathbf{z}) + \log p(\mathbf{z}) - \log q(\mathbf{z}|\mathbf{x}_O)] \\ &\equiv \mathcal{L}_{\text{partial}}. \end{aligned}$$

This bound,  $\mathcal{L}_{\text{partial}}$ , depends only on the observed variables  $\mathbf{x}_O$ , whose dimensionality may vary among different data points. We thus call the inference net,  $q(\mathbf{z}|\mathbf{x}_O)$ , the *partial inference net*. Specifying  $q(\mathbf{z}|\mathbf{x}_O)$  requires distributions for any partition  $\{O, U\}$  of  $I$ .

**Amortized Inference with partial observations.** Inference under partial observations requires the inference net of VAE to be capable to handle arbitrary set of observed data, and sharing parameters across these different sized sets of observations for amortization.

Inspired by the *Point Net (PN)* approach for point cloud classification (Qi et al., 2017; Zaheer et al., 2017), we specify the approximate distribution  $q(\mathbf{z}|\mathbf{x}_O)$  by a *permutation*

*invariant set function encoding*, given by:

$$\mathbf{c}(\mathbf{x}_O) := g(h(\mathbf{s}_1), h(\mathbf{s}_2), \dots, h(\mathbf{s}_{|O|})), \quad (5)$$

where  $\mathbf{s}_d$  carries the information of the input of the  $d$ -th observed variable, and  $|O|$  is the number of observed variables. In particular,  $\mathbf{s}_d$  contains the information about the identity of the input  $\mathbf{e}_d$  and the corresponding input value  $x_d$ . There are many ways to define the identity variable,  $\mathbf{e}_d$ . Naively, it could be the coordinates of observed pixels for images, and one-hot embedding of the number of questions in a questionnaire. With different problem settings, it can be beneficial to learn  $\mathbf{e}$  as an embedding of the identity of the variable, either with or without an naive encoding as input. In this work, we treat  $\mathbf{e}$  as an unknown embedding, to be optimized during training.

There are also different ways to construct  $\mathbf{s}_d$ . A common choice is concatenation,  $\mathbf{s}_d = [\mathbf{e}_d, x_d]$ , which is often used in computer vision applications (Qi et al., 2017). Such architecture is illustrated in Figure 1(a). We refer to this setting as the *Pointnet (PN)* specification of Partial VAE. However, the construction of  $\mathbf{s}_d$  can be more flexible. We propose to construct  $\mathbf{s}_d = \mathbf{e}_d * x_d$  using element-wise multiplication as an alternative, shown in Figure 1(b). We show that this formulation generalizes naive Zero Imputation (ZI) VAE (Nazabal et al., 2018) (cf. Appendix C.1). We refer to the multiplication setting as the *Pointnet Plus (PNP)* specification of Partial VAE.

We can then use a neural network  $h(\cdot)$  to map the input  $\mathbf{s}_d$  to  $\mathbb{R}^K$ , where  $K$  is the latent space size. The key to the PNP/PN structure is the permutation invariant aggregation operation  $g(\cdot)$ , such as max-pooling or summation. In this way, the mapping  $\mathbf{c}(\mathbf{x}_O)$  is invariant to the permutations of elements of  $\mathbf{x}_O$ , and  $\mathbf{x}_O$  can have arbitrary length. Finally, the fixed-size code  $\mathbf{c}(\mathbf{x}_O)$  is fed into an ordinary neural network, that transforms the code into the statistics of a multivariate Gaussian distribution to approximate  $p(\mathbf{z}|\mathbf{x}_O)$ . The procedure is illustrated in Figure 1. As discussed before, given  $p(\mathbf{z}|\mathbf{x}_O)$ , we can estimate  $p(\mathbf{x}_U|\mathbf{z})$ .

### 3.3 Efficient Dynamic Discovery of High-value Information

We now cast the active variable selection problem (1) as an adaptive Bayesian experimental design problem, utilizing  $p(\mathbf{x}_U|\mathbf{x}_O)$  inferred by the Partial VAE. Algorithm 1 summarizes the EDDI framework.

**Information Reward.** We designed a variable selection acquisition function in an information theoretic way following Bayesian experimental design (Lindley, 1956; Bernardo, 1979). Lindley (1956) provides a generic formulation of Bayesian experimental design by maximizing the expected Shannon information. Bernardo (1979) generalizes it by considering the decision task context.



**Algorithm 1** EDDI: Algorithm Overview

---

**Require:** Training dataset  $\mathbf{X}$ , which is partially observed;  
 Test dataset  $\mathbf{X}^*$  with no observations collected yet; Indices  $\phi$  of target variables.

- 1: **Train Partial VAE** by optimizing partial variational bound with  $\mathbf{X}$  (cf. Section 3.2)
- 2: **Actively acquire feature value**  $x_i$  to estimate  $\mathbf{x}_\phi^*$  for each test point (cf. Section 3.3)

**for** each test instance **do**  
      $\mathbf{x}_O \leftarrow \emptyset$  (no variable value has been observed for any test point)  
     **repeat**  
         Choose variable  $x_i$  from  $U \setminus \phi$  to maximize the information reward (Equation (9))  
          $\mathbf{x}_O \leftarrow x_i \cup \mathbf{x}_O$   
     **until** Stopping criterion reached (e.g. the time budget)  
**end for**

---

For a given task, we are interested in statistics of some variables  $\mathbf{x}_\phi$ , where  $\mathbf{x}_\phi \subset \mathbf{x}_U$ . Given a new instance (user), assume that we have observed  $\mathbf{x}_O$  so far for this instance, and we need to select the next variable  $x_i$  (an element of  $\mathbf{x}_{U \setminus \phi}$ ) to observe. Following Bernardo (1979), we select  $x_i$  by maximizing:

$$R(i, \mathbf{x}_O) = \mathbb{E}_{\mathbf{x}_i \sim p(\mathbf{x}_i | \mathbf{x}_O)} D_{KL} [p(\mathbf{x}_\phi | \mathbf{x}_i, \mathbf{x}_O) \| p(\mathbf{x}_\phi | \mathbf{x}_O)]. \quad (6)$$

In our paper, we mainly consider the case that a subset of interesting observations represents the statistics of interest  $\mathbf{x}_\phi$ . Sampling  $\mathbf{x}_i \sim p(\mathbf{x}_i | \mathbf{x}_O)$  is approximated by  $\mathbf{x}_i \sim \hat{p}(\mathbf{x}_i | \mathbf{x}_O)$ , where  $\hat{p}(\mathbf{x}_i | \mathbf{x}_O)$  can be obtained by using the Partial VAE. It is implemented by first sampling  $\mathbf{z} \sim q(\mathbf{z} | \mathbf{x}_O)$ , and then  $\mathbf{x}_i \sim p(\mathbf{x}_i | \mathbf{z})$ . The same applies for  $p(\mathbf{x}_i, \mathbf{x}_\phi | \mathbf{x}_O)$  which appears in Equation (8).

**Efficient approximation of the Information reward.** The Partial VAE allows us to sample  $\mathbf{x}_i \sim p(\mathbf{x}_i | \mathbf{x}_O)$ . However, the KL term in Equation (6),

$$\begin{aligned} D_{KL} [p(\mathbf{x}_\phi | \mathbf{x}_i, \mathbf{x}_O) \| p(\mathbf{x}_\phi | \mathbf{x}_O)] \\ = - \int_{\mathbf{x}_\phi} p(\mathbf{x}_\phi | \mathbf{x}_i, \mathbf{x}_O) \log \frac{p(\mathbf{x}_\phi | \mathbf{x}_O)}{p(\mathbf{x}_\phi | \mathbf{x}_i, \mathbf{x}_O)}, \end{aligned} \quad (7)$$

is intractable since both  $p(\mathbf{x}_\phi | \mathbf{x}_i, \mathbf{x}_O)$  and  $p(\mathbf{x}_\phi | \mathbf{x}_O)$  are intractable. For high dimensional  $\mathbf{x}_\phi$ , entropy estimation could be difficult. The entropy term  $\int_{\mathbf{x}_\phi} p(\mathbf{x}_\phi | \mathbf{x}_i, \mathbf{x}_O) \log p(\mathbf{x}_\phi | \mathbf{x}_i, \mathbf{x}_O)$  depends on  $i$  hence cannot be ignored. In the following, we show how to approximate this expression.

Note that analytic solutions of KL-divergences are available under specific variational distribution families of  $q(\mathbf{z} | \mathbf{x}_O)$  (such as the Gaussian distribution commonly used in VAEs). Instead of calculating the information reward in  $\mathbf{x}$  space,

we have shown that one can compute in the  $\mathbf{z}$  space (cf. Appendix A.1):

$$\begin{aligned} R(i, \mathbf{x}_O) = & \mathbb{E}_{\mathbf{x}_i \sim p(\mathbf{x}_i | \mathbf{x}_O)} D_{KL} [p(\mathbf{z} | \mathbf{x}_i, \mathbf{x}_O) \| p(\mathbf{z} | \mathbf{x}_O)] - \\ & \mathbb{E}_{\mathbf{x}_\phi, \mathbf{x}_i \sim p(\mathbf{x}_\phi, \mathbf{x}_i | \mathbf{x}_O)} D_{KL} [p(\mathbf{z} | \mathbf{x}_\phi, \mathbf{x}_i, \mathbf{x}_O) \| p(\mathbf{z} | \mathbf{x}_\phi, \mathbf{x}_O)]. \end{aligned} \quad (8)$$

Note that Equation (8) is exact. Additionally, we use the partial VAE approximation  $p(\mathbf{z} | \mathbf{x}_\phi, \mathbf{x}_i, \mathbf{x}_O) \approx q(\mathbf{z} | \mathbf{x}_\phi, \mathbf{x}_i, \mathbf{x}_O)$ ,  $p(\mathbf{z} | \mathbf{x}_O) \approx q(\mathbf{z} | \mathbf{x}_O)$  and  $p(\mathbf{z} | \mathbf{x}_i, \mathbf{x}_O) \approx q(\mathbf{z} | \mathbf{x}_i, \mathbf{x}_O)$ . This leads to the final approximation of the information reward:

$$\begin{aligned} \hat{R}(i, \mathbf{x}_O) = & \mathbb{E}_{\mathbf{x}_i \sim \hat{p}(\mathbf{x}_i | \mathbf{x}_O)} D_{KL} [q(\mathbf{z} | \mathbf{x}_i, \mathbf{x}_O) \| q(\mathbf{z} | \mathbf{x}_O)] - \\ & \mathbb{E}_{\mathbf{x}_\phi, \mathbf{x}_i \sim \hat{p}(\mathbf{x}_\phi, \mathbf{x}_i | \mathbf{x}_O)} D_{KL} [q(\mathbf{z} | \mathbf{x}_\phi, \mathbf{x}_i, \mathbf{x}_O) \| q(\mathbf{z} | \mathbf{x}_\phi, \mathbf{x}_O)]. \end{aligned} \quad (9)$$

With this approximation, the divergence between  $q(\mathbf{z} | \mathbf{x}_i, \mathbf{x}_O)$  and  $q(\mathbf{z} | \mathbf{x}_O)$  can often be computed analytically in the Partial VAE setting, for example, under Gaussian parameterization. The only Monte Carlo sampling required is the one set of samples  $\mathbf{x}_\phi, \mathbf{x}_i \sim p(\mathbf{x}_\phi, \mathbf{x}_i | \mathbf{x}_O)$  that can be shared across different KL terms in Equation (9). Our EDDI framework is opensource at <https://github.com/Microsoft/EDDI>.

## 4 Experiments

Here we evaluate the proposed EDDI framework. We first assess the Partial VAE component of EDDI alone on an image inpainting task both qualitatively and quantitatively (Section 4.1). We compare our proposed two PN-based Partial VAE with the zero-imputing (ZI) VAE (Nazabal et al., 2018). Additionally, we modify the ZI VAE to use the mask matrix indicating which variables are currently observed as input. We name this method ZI-m VAE. We then demonstrate the performance of the entire EDDI framework on datasets from the UCI repository (Section 4.2), as well as in two real-life application scenarios: Risk assessment in intensive care (Section 4.3) and public health assessment with national health survey (Section 4.4). We compare the performance of EDDI, using four different Partial VAE settings, with three baseline information acquisition strategies. The first baseline is the *random active feature selection strategy* (denoted as *RAND*) which randomly picks the next variable to observe. *RAND* reflects the strategy used in many real-world applications, such as online surveys. The second baseline method is the *single best strategy* (denoted as *SING*) which finds a single fixed global optimal order of selecting variables. This order is then applied to all data points. *SING* uses the objective function as in Equation (9) to find the optimal ordering by averaging over all the test data.

### 4.1 Image inpainting with Partial VAE

We evaluate the performance of Partial VAE with the image inpainting task, which is to fill in the removed pixels. We




Figure 2: Image inpainting example with MNIST dataset using Partial VAE with four settings.




Figure 3: Information curves of active variable selection, demonstrated on three UCI datasets (based on PNP parameterization of Partial VAE). This displays negative test RMSE (y axis, the lower the better) during the course of active selection (x-axis). Error bars represent standard errors over 10 runs.

perform the evaluation in two different settings: in the first setting, pixels are randomly removed, and in the second setting, a continuous patch of pixels are removed.

Table 1: Comparing models trained on partially observed MNIST. VAE-full is an ideal reference.

Method	VAE-full	ZI	ZI-m	PN	PNP
Train ELBO	-95.05	<b>-113.64</b>	-117.29	-121.43	<b>-113.64</b>
Test ELBO (Rnd.)	-101.46	-116.01	-118.61	-122.20	<b>-114.01</b>
Test ELBO (Reg.)	-101.46	-130.61	-123.87	-116.53	<b>-113.19</b>

**Inpainting Random Missing Pixels.** We use MNIST dataset (LeCun, 1998) and remove pixels randomly for this task. The same settings are used for all methods (see Appendix B.1 for details). During training, we remove a random portion (uniformly sampled between 0% and 70%) of pixels. We then impute missing pixels on a partially observed test set (constructed by removing 70% of the pixels uniform randomly). The performance of pixel imputation is evaluated by test ELBOs on missing pixels. The first two rows in Table 1 show training and test ELBOs for all algorithms using this partially observed dataset. Addition-

ally, we show ordinary VAE (VAE-full) trained on the fully observed dataset as an ideal reference. Among all Partial VAE methods, the PNP approach performs best.

**Inpainting Regions.** We then consider inpainting large contiguous regions of images. It aims to evaluate the capability of the Partial VAEs to produce all possible outcomes with better uncertainty estimates. With the same trained model as before, we remove the region of the upper 60% pixels of the image in the test set. We then evaluate the average likelihoods of the models. The last row of Table 1 shows the results of the test ELBO in this case. PNP based Partial VAE performs better than other settings. Note that given only the lower half of a digit, the number cannot be identified uniquely. ZI (Figure 2(b)) fails to cover the different possible modes due to its limitation in posterior inference. ZI-m (Figure 2(c)) is capable of producing multiple modes. However, some of the generated samples are not consistent with the given part (i.e., some digits of 2 are generated). Our proposed PN (Figure 2(d)) and PNP (Figure 2(e)) are capable of recovering different modes, and are consistent with observations.




Figure 4: First four decision steps on Boston Housing test data. EDDI is “personalized” comparing SING. Full names of the variables are listed in the Appendix B.2.




Figure 5: Comparison of DRAL (Lewenberg et al., 2017) and EDDI on Boston Housing dataset. EDDI out performs DRAL significantly regarding test RMSE in every step.

Method	Time
DRAL	2747.16
EDDI	<b>2.64</b>

Table 2: Test CPU time (in seconds) per test point for active variable selection using EDDI and DRAL. EDDI is  $10^3$  times more efficient than DRAL (Lewenberg et al., 2017) computationally.

## 4.2 EDDI on UCI datasets

Table 3: Average ranking of AUIC over 6 UCI datasets.

Method	ZI	ZI-m	PNP	PN
EDDI	5.72 (0.03)	5.54 (0.02)	<b>5.08 (0.02)</b>	5.25 (0.02)
Random	8.03 (0.03)	8.10 (0.03)	7.77 (0.03)	7.79 (0.03)
Single best	8.68 (0.03)	5.50 (0.02)	5.20 (0.02)	5.28 (0.02)

Given the effectiveness of our proposed Partial VAE, we now demonstrate the performance of our proposed EDDI framework in comparison with random selection (RAND) and single optimal ordering (SING). We first apply EDDI on 6 different UCI datasets (cf. Appendix B.2) (Dheeru & Karra Taniskidou, 2017). We report the results of EDDI with all these four different specifications of Partial VAE (ZI, ZI-m, PN, PNP).

All Partial VAE are first trained on partially observed UCI datasets where a random portion of variables is removed. We actively select variable for each test point starting with empty observation  $\mathbf{x}_o = \emptyset$ . In all UCI datasets, we randomly sample 10% of the data as the test set. All experiments are repeated for ten times.

Taking PNP based setting as an example, Figure 3 shows the test RMSE on  $\mathbf{x}_\phi$  for each variable selection step with three different datasets, where  $\mathbf{x}_\phi$  is defined by the UCI task. We call this curve the *information curve (IC)*. We see that EDDI can obtain information efficiently. It archives the same test RMSE with less than half of the variables. Single optimal ordering also improves upon random ordering. However, it is less efficient compared with EDDI, since EDDI perform active learning for each data instance which is “personalized”. Figure 4 shows an example of the decision processes using EDDI and SING. The first step of EDDI overlaps largely with SING. From the second step, EDDI makes “personalized” decisions.

We also present the average performance among all datasets with different settings. The area under the information curve (AUIC), can then be used to compare the performance across models and strategies. Smaller AUIC value indicates better performance. However, due to different

datasets have different scales of RMSEs and different numbers of variables (indicated by steps), it is not fair to average the AUIC across datasets to compare overall performances. We thus define average *ranking* of AUIC that compares 12 methods (indexed by  $i$ ) averaging these datasets as:  $r_i = \frac{1}{\sum_j N_j} \sum_{j=1}^6 \sum_{k=1}^{N_j} r_{ijk}$ ,  $i = 1, \dots, 12$ . These 12 methods are cross combinations of four Partial VAE models with three variable selection strategies.  $r_i$  is the final ranking of  $i$ th combination,  $r_{ijk}$  is the ranking of the  $i$ th combination (based on AUIC value) regarding the  $k$ th test data point in the  $j$ th UCI dataset, and  $N_j$  is the size of the  $j$ th UCI dataset. This gives us  $6 \sum_j N_j$  different rankings. Finally, we compute the mean and standard error statistics based on these rankings. Table 3 summarize the average ranking results. We provide additional statistical significance test (Wilcoxon signed-rank test for paired data) in Appendix B.2.2. Based on these experimental results, we see that EDDI outperforms other variable selection order in all different Partial VAE settings. Among different partial VAE settings, PNP/PN-based settings perform better than ZI-based settings.

**Comparison with non-amortized method.** Additionally, we compare EDDI to DRAL (Lewenberg et al., 2017) which is the state-of-the-art method for the same problem setting. As discussed in Section 2, DRAL is linear and requires high computational cost. The DRAL paper only tested their method on a single test data point due to its limitation on computational efficiency. We compare DRAL with EDDI on Boston Housing dataset with ten randomly selected test points here. Results are shown in Figure 5, where EDDI significantly outperforms DARE thanks to more flexible Partial VAE model. Additionally, EDDI is 1000 times more efficient than DARE as shown in Table 2.

## 4.3 Risk assessment with MIMIC-III

We now apply EDDI to risk assessment tasks using the Medical Information Mart for Intensive Care (MIMIC III) database (Johnson et al., 2016). MIMIC III is the most extensive publicly available clinical database, containing real-world records from over 40,000 critical care patients




Figure 6: Information curves (based on Bernoulli negative log likelihood) of active variable selection on risk assessment task on MIMIC III with PNP setting.




Figure 7: Information curves of active (grouped) variable selection on risk assessment task on NHANES with PNP setting.

with 60,000 ICU stays. The risk assessment task is to predict the final mortality. We preprocess the data for this task following Harutyunyan et al. (2017)<sup>1</sup>. This results in a dataset of 21139 patients. We treat the final mortality of a patient as a Bernoulli variable. For our task, we focus on variable selection, which corresponds to medical instrument selection. We thus further process the time series variables into static variables based on temporal averaging.

Figure 6 shows the information curve (based on Bernoulli likelihoods) of different strategies, using PNP based Partial VAE as an example (more results in Appendix B.3). Table 4 shows the average ranking of AUIC with different settings. In this application, EDDI significantly outperforms other variable selection strategies in all different settings of Partial VAE, and PNP based setting performs best.

#### 4.4 Public Health Assessment with NHANES

Finally, we apply our methods to public health assessment using NHANES 2015-2016 data (cdc, 2005). NHANES is a program with adaptable components of measurements, to assess the health and nutritional status of adults and children in the United States. Every year, thousands individuals of all ages are interviewed and examined in their homes. This 2015-2016 NHANES data contains three major sections, the questionnaire interview, examinations and lab tests for 9971 subjects in the publicly available version of this cycle. In our setting, we consider the whole set of lab test results (139 variables) as the target variable of interest  $\mathbf{x}_\phi$  since they are expensive and reflects the subject’s health status, and we active select the questions from the extensive questionnaire

<sup>1</sup><https://github.com/yerevann/mimic3-benchmarks>

Method	EDDI	Random	Single best
ZI	8.83 (0.01)	7.97 (0.02)	9.83 (0.01)
ZI-m	4.91 (0.01)	7.00 (0.01)	5.91 (0.01)
PN	4.96 (0.01)	6.62 (0.01)	5.96 (0.01)
PNP	<b>4.39 (0.01)</b>	6.18 (0.01)	5.39 (0.01)

Table 4: Average ranking on AUIC of MIMIC III

Method	EDDI	Random	Single best
ZI	6.00 (0.10)	8.45 (0.09)	6.51 (0.09)
ZI-m	8.06 (0.09)	8.67 (0.09)	8.68 (0.07)
PN	5.28 (0.10)	5.57 (0.10)	5.46 (0.09)
PNP	<b>4.80 (0.10)</b>	5.30 (0.10)	5.17 (0.10)

Table 5: Average ranking on AUIC of NHANES

(665 variables).

The questionnaire of NHANES is divided into 73 different groups. In practice, questions in the same group are often examined together. Therefore, we perform active variable selection on the group level: at each step, the algorithm selects one group to observe. This is more challenging than the experiments in previous sections since it requires the generative model to simulate a group of unobserved data in Equation (9) at the same time. When evaluating test RMSE on the target variable of interest, we treat variables in each group equally. For a fair comparison, the calculation of the area under the information curve (AUIC) is weighted by the size of the group chosen by the algorithms. Specifically, AUIC is calculated after spline interpolation. The information curve plots in Figure 7, together with Table 5 of AUIC statistics show that our EDDI outperforms other baselines. In addition, this experiment shows that EDDI is capable of performing active selection on a large pool of grouped variables to estimate a high dimensional target.

## 5 Conclusion

In this paper, we present EDDI, a novel and efficient framework for dynamic active variable selection for each instance. Within the EDDI framework, we propose Partial VAE which performs amortized inference to handle missing data. Partial VAE alone can be used as a non-linear computational efficient probabilistic imputation method. Based on it, we design a variable wise acquisition function for EDDI and derive corresponding approximation method. EDDI has demonstrated its effectiveness on active variable selection tasks across multiple real-world applications. In the future, we would extend the EDDI framework to handle more complicated scenarios, such as data missing not at random, time-series, and the cold-start situation.



## References

- National health and nutrition examination survey, 2005. URL <https://www.cdc.gov/nchs/nhanes/>.
- Bernardo, J. M. Expected information as expected utility. *The Annals of Statistics*, pp. 686–690, 1979.
- Blei, D. M., Ng, A. Y., and Jordan, M. I. Latent dirichlet allocation. *Journal of Machine Learning Research*, 3 (Jan):993–1022, 2003.
- Dempster, A. P., Laird, N. M., and Rubin, D. B. Maximum likelihood from incomplete data via the em algorithm. *Journal of the Royal Statistical Society. Series B (methodological)*, pp. 1–38, 1977.
- Dheeru, D. and Karra Taniskidou, E. UCI machine learning repository, 2017. URL <http://archive.ics.uci.edu/ml>.
- Garnelo, M., Rosenbaum, D., Maddison, C. J., Ramalho, T., Saxton, D., Shanahan, M., Teh, Y. W., Rezende, D. J., and Eslami, S. Conditional neural processes. *arXiv preprint arXiv:1807.01613*, 2018.
- Gopalan, P. K., Charlin, L., and Blei, D. Content-based recommendations with poisson factorization. In *Advances in Neural Information Processing Systems*, pp. 3176–3184, 2014.
- Hamesse, C., Ackermann, P., Kjellström, H., and Zhang, C. Simultaneous measurement imputation and outcome prediction for achilles tendon rupture rehabilitation. In *ICML/IJCAI Joint Workshop on Artificial Intelligence in Health*, 2018.
- Harutyunyan, H., Khachatrian, H., Kale, D. C., and Galstyan, A. Multitask learning and benchmarking with clinical time series data. *arXiv preprint arXiv:1703.07771*, 2017.
- Houlsby, N., Huszár, F., Ghahramani, Z., and Lengyel, M. Bayesian active learning for classification and preference learning. *arXiv preprint arXiv:1112.5745*, 2011.
- Huang, S.-J., Xu, M., Xie, M.-K., Sugiyama, M., Niu, G., and Chen, S. Active feature acquisition with supervised matrix completion. *arXiv preprint arXiv:1802.05380*, 2018.
- Jain, P., Meka, R., and Dhillon, I. S. Guaranteed rank minimization via singular value projection. In *Advances in Neural Information Processing Systems*, 2010.
- Johnson, A. E., Pollard, T. J., Shen, L., Li-wei, H. L., Feng, M., Ghassemi, M., Moody, B., Szolovits, P., Celi, L. A., and Mark, R. G. Mimic-iii, a freely accessible critical care database. *Scientific Data*, 3:160035, 2016.
- Keshavan, R. H., Montanari, A., and Oh, S. Matrix completion from noisy entries. *Journal of Machine Learning Research*, 2010.
- Kingma, D. P. and Ba, J. L. Adam: a method for stochastic optimization. In *International Conference on Learning Representations*, pp. 1–13, 2015.
- Kingma, D. P. and Welling, M. Auto-encoding variational bayes. In *International Conference on Learning Representation*, 2014.
- LeCun, Y. The mnist database of handwritten digits. <http://yann.lecun.com/exdb/mnist/>, 1998.
- Lewenberg, Y., Bachrach, Y., Paquet, U., and Rosenschein, J. S. Knowing what to ask: A bayesian active learning approach to the surveying problem. In *AAAI*, pp. 1396–1402, 2017.
- Lindley, D. V. On a measure of the information provided by an experiment. *The Annals of Mathematical Statistics*, pp. 986–1005, 1956.
- Little, R. and Rubin, D. Statistical analysis with missing data. Technical report, J. Wiley, 1987.
- MacKay, D. J. Information-based objective functions for active data selection. *Neural computation*, 4(4):590–604, 1992.
- McCallumzy, A. K. and Nigamy, K. Employing em and pool-based active learning for text classification. In *International Conference on Machine Learning*, pp. 359–367. Citeseer, 1998.
- Melville, P., Saar-Tsechansky, M., Provost, F., and Mooney, R. Active feature-value acquisition for classifier induction. In *International Conference on Data Mining*, pp. 483–486. IEEE, 2004.
- Nazabal, A., Olmos, P. M., Ghahramani, Z., and Valera, I. Handling incomplete heterogeneous data using vaes. *arXiv preprint arXiv:1807.03653*, 2018.
- Qi, C. R., Su, H., Mo, K., and Guibas, L. J. Pointnet: Deep learning on point sets for 3d classification and segmentation. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pp. 652–660, 2017.
- Ranganath, R., Tran, D., and Blei, D. Hierarchical variational models. In *International Conference on Machine Learning*, pp. 324–333, 2016.
- Rezende, D. J., Mohamed, S., and Wierstra, D. Stochastic backpropagation and approximate inference in deep generative models. In *International Conference on Machine Learning*, 2014.

- Rubin, D. B. Inference and missing data. *Biometrika*, 63(3): 581–592, 1976.
- Saar-Tsechansky, M., Melville, P., and Provost, F. Active feature-value acquisition. *Management Science*, 55(4): 664–684, 2009.
- Salakhutdinov, R. and Mnih, A. Bayesian probabilistic matrix factorization using markov chain monte carlo. In *International conference on Machine learning*, pp. 880–887. ACM, 2008.
- Scheffer, J. Dealing with missing data. 2002.
- Settles, B. Active learning. *Synthesis Lectures on Artificial Intelligence and Machine Learning*, 6(1):1–114, 2012.
- Shim, H., Hwang, S. J., and Yang, E. Joint active feature acquisition and classification with variable-size set encoding. In *Advances in Neural Information Processing Systems*, 2018.
- Stern, D., Herbrich, R., and Graepel, T. Matchbox: Large scale bayesian recommendations. In *International World Wide Web Conference*, 2009.
- Thahir, M., Sharma, T., and Ganapathiraju, M. K. An efficient heuristic method for active feature acquisition and its application to protein-protein interaction prediction. In *BMC proceedings*, volume 6, pp. S2. BioMed Central, 2012.
- Wang, C. and Blei, D. M. Collaborative topic modeling for recommending scientific articles. In *International Conference on Knowledge Discovery and Data Mining*, pp. 448–456. ACM, 2011.
- Wu, G., Domke, J., and Sanner, S. Conditional inference in pre-trained variational autoencoders via cross-coding. *arXiv preprint arXiv:1805.07785*, 2018.
- Yu, H.-F., Rao, N., and Dhillon, I. S. Temporal regularized matrix factorization for high-dimensional time series prediction. In *Advances in Neural Information Processing Systems*, pp. 847–855, 2016.
- Zaheer, M., Kottur, S., Ravanbakhsh, S., Poczos, B., Salakhutdinov, R. R., and Smola, A. J. Deep sets. In *Advances in Neural Information Processing Systems*, pp. 3394–3404, 2017.
- Zakim, D., Braun, N., Fritz, P., and Alscher, M. D. Underutilization of information and knowledge in everyday medical practice: Evaluation of a computer-based solution. *BMC Medical Informatics and Decision Making*, 8(1):50, 2008.
- Zhang, C., Butepage, J., Kjellstrom, H., and Mandt, S. Advances in variational inference. *arXiv preprint arXiv:1711.05597*, 2017.
- Zheng, Z. and Padmanabhan, B. On active learning for data acquisition. In *International Conference on Data Mining*, pp. 562–569. IEEE, 2002.